

Time-series, PostgreSQL, and You

Mathis Van Eetvelde



About me



Mathis Van Eetvelde

Developer Advocate
@Timescale



@meetvelde



@mathisve



Mathis Van Eetvelde



mathisvaneetvelde.com

Agenda

Topics we will cover today

- 01** What is time-series data
- 02** Time-series data on PostgreSQL
- 03** Time-series data and You
- 03.1** TimescaleDB



01

What is time-series data?



Time-series DDL

Financial time-series data

```
CREATE TABLE stocks (  
  time TIMESTAMPTZ NOT NULL,  
  price INT NULL,  
  symbol TEXT NULL  
);
```

IoT time-series data

```
CREATE TABLE measurements (  
  time TIMESTAMPTZ NOT NULL,  
  deviceID INT NULL,  
  sensor_value REAL NULL,  
  location geometry NULL  
);
```

...

Kinds of time-series data

- Financial data
 - Stocks
 - OLTP (Online Transaction Processing)
- IoT data
 - Smartphones
 - Smart fridges, TVs, etc.
 - Industrial sensors
- Observability data
 - Traces, spans
 - Logs
 - Metrics





Time-series data means scale

NYSE transactions per **day**

4.2B

Consolidated Tape A
trading volume
<https://www.nyse.com/markets/us-equity-volumes>

NYSE transactions per **second**

162K

4.2B Transactions

23400 Trading Seconds

NYSE transactions in **5 years**

5.4T

CFTC requires that financial
records be kept for 5 years
and “readily available” for
2 years

02

Time-series data on PostgreSQL

Why PostgreSQL?



- Flourishing ecosystem
- Tried and tested
- Not a DSL (InfluxQL, FQL, etc.)
- You don't have to write your own JOINS on the application level
- You already know PostgreSQL




Hard limitations of PostgreSQL


- 4.2 Billion databases
- 1.4 Billion tables per database
- 32 TB per table (8192 bytes block size)
- 4.2 Billion pages per table



Partitioning



Partitioning



Partitioning

- Smaller tables
- Smaller indexes
- Partition Pruning



Partitioning

- Inheritance based partitioning (pre PG 10)
 - Harder
 - Versatile
- Declarative partitioning since (PG 10)
 - Easier





Inheritance based partitioning

Create the root table

```
CREATE TABLE measurement (  
    time TIMESTAMPTZ NOT NULL,  
    sensorID INT NULL,  
    sensor_value REAL NULL  
);
```

Create inheritance tables

```
CREATE TABLE measurement_y2022m01 () INHERITS (measurement);  
CREATE TABLE measurement_y2022m02 () INHERITS (measurement);  
CREATE TABLE measurement_y2022m03 () INHERITS (measurement);
```

Create indexes

```
CREATE INDEX ON measurement_y2022m01 (time);  
CREATE INDEX ON measurement_y2022m02 (time);  
CREATE INDEX ON measurement_y2022m03 (time);
```



Trigger function

```
CREATE OR REPLACE FUNCTION measurement_insert_trigger()
RETURNS TRIGGER AS $$
BEGIN
    IF ( NEW.time >= DATE '2022-01-01' AND
        NEW.time < DATE '2022-02-01' ) THEN
        INSERT INTO measurement_y2022m01 VALUES (NEW.*);

    ...
    ELSIF ( NEW.time >= DATE '2022-03-01' AND
        NEW.time < DATE '2022-04-01' ) THEN
        INSERT INTO measurement_y2022m03 VALUES (NEW.*);
    ELSE
        RAISE EXCEPTION 'Date out of range.';
    END IF;
    RETURN NULL;
END;
$$
LANGUAGE plpgsql;
```

```
CREATE TRIGGER insert_measurement_trigger
BEFORE INSERT ON measurement
FOR EACH ROW EXECUTE FUNCTION measurement_insert_trigger();
```




Inserts

```
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-01-3', 2, 2.9);
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-01-6', 1, 3.5);
...
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-03-27', 3, 2.5);
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-03-29', 2, 5.3);
```

```
SELECT tableoid, * from measurement;
```

tableoid	time	sensorid	sensor_value
19103	2022-01-03 00:00:00+00	2	2.9
19103	2022-01-06 00:00:00+00	1	3.5
...
19109	2022-03-27 00:00:00+00	3	2.5
19109	2022-03-29 00:00:00+00	2	5.3



Partition Pruning - on

```
SET enable_partition_pruning = on; -- default
```

```
EXPLAIN ANALYZE SELECT * FROM measurement WHERE time > '2022-03-01';
```

```
Index Scan using measurementy2022m03_time_idx on measurementy2022m03 measurement  
(cost=0.15..23.05 rows=617 width=16) (actual time=0.005..0.006 rows=5 loops=1)
```

```
Index Cond: ("time" > '2022-03-01 00:00:00+00'::timestamp with time zone)
```

```
Planning Time: 0.090 ms
```

```
Execution Time: 0.019 ms
```

```
(4 rows)
```



Partition Pruning - off

```
SET enable_partition_pruning = off;
```

```
EXPLAIN ANALYZE SELECT * FROM measurement WHERE time > '2022-03-01';
```

```
Append (cost=0.15..78.41 rows=1851 width=16) (actual time=0.031..0.033 rows=5 loops=1)
```

```
-> Index Scan using measurementy2022m01_time_idx on measurementy2022m01 measurement_1
```

```
Index Cond: ("time" > '2022-03-01 00:00:00+00'::timestamp with time zone)
```

```
-> Index Scan using measurementy2022m02_time_idx on measurementy2022m02 measurement_2
```

```
Index Cond: ("time" > '2022-03-01 00:00:00+00'::timestamp with time zone)
```

```
-> Index Scan using measurementy2022m03_time_idx on measurementy2022m03 measurement_3
```

```
Index Cond: ("time" > '2022-03-01 00:00:00+00'::timestamp with time zone)
```

```
Planning Time: 0.123 ms
```

```
Execution Time: 0.053 ms
```

```
(9 rows)
```



Declarative Partitioning

Create the root table

```
CREATE TABLE measurement (  
    time TIMESTAMPTZ NOT NULL,  
    sensorID INT NULL,  
    sensor_value REAL NULL  
) PARTITION BY RANGE (time);
```

Create partitions

```
CREATE TABLE measurementy2022m01 PARTITION OF measurement  
    FOR VALUES FROM ('2022-01-01') TO ('2022-02-01');  
CREATE TABLE measurementy2022m02 PARTITION OF measurement  
    FOR VALUES FROM ('2022-02-01') TO ('2022-03-01');  
CREATE TABLE measurementy2022m03 PARTITION OF measurement  
    FOR VALUES FROM ('2022-03-01') TO ('2022-04-01');
```



Inserts

```
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-01-3', 2, 2.9);
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-01-6', 1, 3.5);
...
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-03-27', 3, 2.5);
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-03-29', 2, 5.3);
```

```
SELECT tableoid, * from measurement;
```

tableoid	time	sensorid	sensor_value
18918	2022-01-03 00:00:00+00	2	2.9
18918	2022-01-06 00:00:00+00	1	3.5
...
18924	2022-03-27 00:00:00+00	3	2.5
18924	2022-03-29 00:00:00+00	2	5.3



Indexes

```
CREATE INDEX ON measurement (time);
```

```
SELECT  
    tablename, indexname  
FROM  
    pg_indexes  
WHERE  
    tablename LIKE 'measurement%';
```

<i>tablename</i>	<i>indexname</i>
<i>measurement</i>	<i>measurement_time_idx</i>
<i>measurementy2022m01</i>	<i>measurementy2022m01_time_idx</i>
<i>measurementy2022m02</i>	<i>measurementy2022m02_time_idx</i>
<i>measurementy2022m03</i>	<i>measurementy2022m03_time_idx</i>

Declarative partitioning

- Range partitioning
- List partitioning
- Hash partitioning
- Composite partitioning



03

Time-series data and You



Solutions

- Manual labor
- Diy bash script
- pg_partman & pg_cron

TimescaleDB

Postgres for time-series



Quick install

Install using your favorite package manager

```
apt install timescaledb-2-postgresql-14
```

```
yum install timescaledb-2-postgresql-14
```

```
sudo pacman -Syu timescaledb timescaledb-tune
```

Use TimescaleDB in a container

```
docker run -d --name timescaledb -p 5432:5432
```

Or install from source

```
git clone https://github.com/timescale/timescaledb.git
cd timescaledb
git checkout 2.5.1
./bootstrap
cd build && make
make install
```

<https://docs.timescale.com/install/latest/>



TimescaleDB

Create a table

```
CREATE TABLE measurement (  
    time TIMESTAMPTZ NOT NULL,  
    sensorID INT NULL,  
    sensor_value REAL NULL  
);
```

Create a hypertable

```
SELECT create_hypertable('measurement', 'time');
```

Create a distributed hypertable

```
SELECT create_distributed_hypertable('measurement', 'time', 'sensorID');
```





Inserts

```
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-01-3', 2, 2.9);
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-01-6', 1, 3.5);
...
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-03-27', 3, 2.5);
INSERT INTO measurement (time, sensorID, sensor_value) VALUES ('2022-03-29', 2, 5.3);
```

```
SELECT tableoid, * from measurement;
```

tableoid	time	sensorid	sensor_value
19204	2022-01-03 00:00:00+00	2	2.9
19210	2022-01-06 00:00:00+00	1	3.5
19216	2022-01-15 00:00:00+00	3	6.2
19222	2022-01-23 00:00:00+00	1	3.2
...
19240	2022-02-13 00:00:00+00	2	2.6
19246	2022-02-24 00:00:00+00	2	1.7
19246	2022-02-28 00:00:00+00	1	6.1
19252	2022-03-06 00:00:00+00	3	5.2
19252	2022-03-08 00:00:00+00	2	2.3



Auxiliary functions

```
SELECT show_chunks('measurement');
```

```
SELECT drop_chunks('measurement', INTERVAL '5 months');
```

```
SELECT reorder_chunk('_timescaledb_internal._hyper_4_78_chunk', 'measurement_time_idx');
```


```
SELECT attach_tablespace('disk2', 'measurement');
```


```
SELECT move_chunk(  
  chunk => '_timescaledb_internal._hyper_4_78_chunk',  
  destination_tablespace => 'disk2',  
  index_destination_tablespace => 'disk2',  
  reorder_index => 'measurement_time_idx',  
);
```

<https://docs.timescale.com/api/latest/hypertable/>



Continuous aggregates





⌵ **B** (PostgreSQL)

```
select day as time, avg from measurement_daily;
```

Format as

Time series ▾

Query Builder

Show Help >





Continuous aggregates

```
CREATE MATERIALIZED VIEW measurement_daily
WITH (timescaledb.continuous) AS
SELECT
    time_bucket('1 day', "time") AS day,
    avg(sensor_value) AS avg,
    count(sensor_value) AS count
FROM measurement
GROUP BY day;
```

```
SELECT day AS time, avg FROM measurement_daily;
```



Compression

```
ALTER TABLE measurement SET (  
    timescaledb.compress,  
    timescaledb.compress_orderby = 'time DESC',  
    timescaledb.compress_segmentby = 'sensorID'  
);  
  
SELECT add_compression_policy('measurement', INTERVAL '1 month');
```

Up to 94% compression!



Compression

As a simplified example, you might have a table that looks like this to start with:

time	device_id	cpu	energy_consumption
12:00:02	1	88.2	0.8
12:00:02	2	300.5	0.9
12:00:01	1	88.6	0.85
12:00:01	2	299.1	0.95

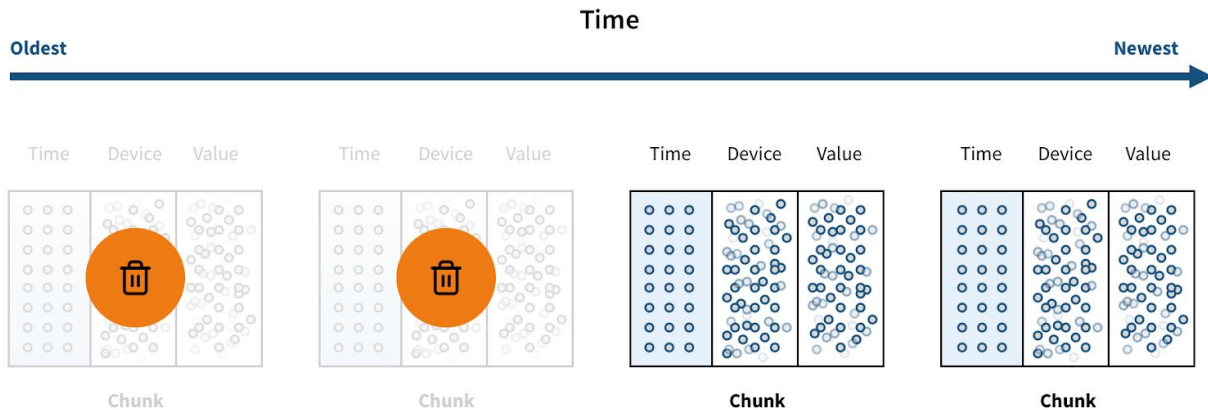
When compression is applied, the data is converted to a single row containing an array, like this:

time	device_id	cpu	energy_consumption
[12:00:02, 12:00:02, 12:00:01, 12:00:01]	[1, 2, 1, 2]	[88.2, 300.5, 88.6, 299.1]	[0.8, 0.9, 0.85, 0.95]



Data Retention

```
SELECT add_retention_policy('measurement', INTERVAL '5 years');
```



Timescale Cloud

TimescaleDB in the cloud



Timescale Cloud

Create a service

Service summary

[Advanced configuration](#)

You can reconfigure your service anytime later on!

Region

us-east-1

High-availability

Enable replicas for your mission critical workloads.

☐

Enable replica

Configuration

Scale later on

0.5 CPU / 2 GB RAM / 10 GB Storage



Create service

Pricing

30-day trial

~~\$0.053~~ \$0 / hour

Compute

\$0.041 \$0 / hour

Storage

\$0.012 \$0 / hour

Monthly est. ⓘ

\$39 \$0 / mo

Interactive demo. Deploy a service with a [demo dataset](#) to learn more about TimescaleDB.



Timescale Cloud

Mathis' databases 

db-62780

SINGLE-NODE

● Running

CPU

2

RAM

8 GB

Disk utilization

77%

7.7 GB of 10 GB

Region

us-east-1

Created a month ago

pgconf nyc demo

SINGLE-NODE

● Running

CPU

2

RAM

8 GB

Disk utilization

3%

3 GB of 100 GB

Region


us-west-2

Created 8 days ago

[+ Create service](#)



Timescale Cloud

db-62780  Running ...

Overview

Explorer


Operations

Metrics

Logs

Settings

General information

Tables 

company

cr2

cr3

cr4

docs

docs_gin


measurement


stocks_real_time


test

textfun

timetest

Continuous aggregates 

measurement 

Size 
4.5 GB

Approx. row count
19.2 M

Number of chunks
53

Chunk time interval
7 days

Schema


Indexes





Foreign Keys

Chunks

Cont Aggs

Policies

 Search

Column name 	Type 	Default 	Nullable 
time	timestamp with time zone		-
sensorid	integer		✓
sensor_value	real		✓



Timescale Cloud

db-62780



Running



Overview

Explorer

Operations

Metrics

Logs

Settings

☒ Enable refresh

Last hour




CPU

0.20%

4 mCPU

14:31




Memory

63.50%

5.08 GB

14:31




Storage

76.92%

7.692 GB

14:31





Timescale Cloud

multinode-test

MULTI-NODE

● Running

Nodes		CPU/node	RAM/node	Region
Access nodes	1	8	32 GB	us-east-1
Data nodes	3	2	8 GB	

Highest disk utilizations

bafww0amr8	1%	<div><div></div></div> 420 MB of 50 GB	nk15jlm6nm	0%	<div><div></div></div> 340 MB of 200 GB
fyod6yv530	0%	<div><div></div></div> 340 MB of 200 GB	m0scp73hko	0%	<div><div></div></div> 340 MB of 200 GB

Questions?



Scan the QR code to visit the
Timescale home page!



Thank you!



Scan the QR code to visit the
Timescale home page!

